

1 Software Architecture

The solution is divided into four main components:

- The **client** - responsible for interfacing with the server or other methods of data input
- The **pathfinding** module - handling graph search and pathing between two points in any space
- The **simulation** module - handling the route selection and the actual "drone control" with the use of pathfinding
- The **visualisation** module - responsible for generating visualisations of the output from the simulation module

Most of the inter-module design choices explained below, are fueled by the dependency inversion principle and the rest by good general OOP practice. Due to the size of the solution, a lot of the simpler classes are missed out of the diagrams. The descriptions below will revolve mostly around the key decisions behind the architecture and their benefits/alternatives.

1.1 Client module

The first choice is deciding how input data is to be parsed and prepared for consumption in other modules. Three "Data" classes were appointed for this role. These closely match the structure of the json data provided, yet they do not copy the structure entirely. The alternatives included entirely omitting intermediate classes such as these, however should any change in structure of data happen, changes in all the classes consuming the data would need to be made, or at least their parsing code, in this case, we simply change the parsing code in the ClientService which is responsible for generating the Data classes from the input data. The ClientService is modelled as an interface, in the case that the entire format of the data changes, a new ClientService implementation can be written to handle it. Figure 1 shows the class diagram of the client module. In the case of this problem, the ClientService needs to communicate with a server and so the [AQWebServerClient](#) is modelled as an example. The [AQSensor](#) class is shown to show the relationship between the Client and the Simulation module (discussed in detail below), following dependency inversion principles, defining the concrete sensor class in this module, decouples the input data completely from the classes consuming it.

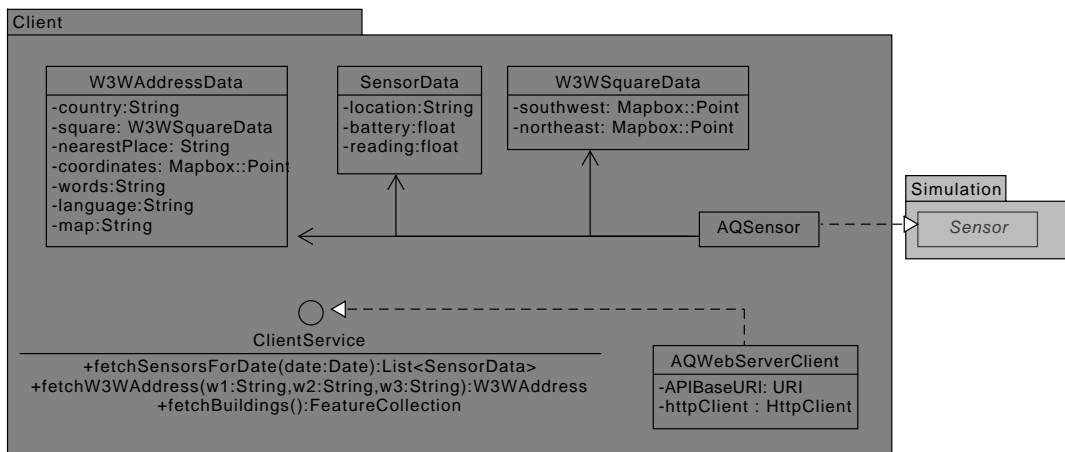


Figure 1: UML diagram of the client module

1.2 Pathfinding module

A big part of the problem is finding a path between two points on a plane, this is a widely studied area and a problem which can be easily abstracted away from the specific domain of drone sensor data collection. To represent the abstract domain, the Graph interface and [SearchNode](#) abstract class are created, each Graph provides at its core a method to retrieve the neighbours of each [SearchNode](#), each of which represents a point and path needing to be taken to reach that point. This representation allows the creation of algorithms able to find paths in a multitude of different domains. An alternative choice was skipping those classes and embedding the graph of the domain in each class modelling a pathfinding algorithm, this approach however would force the creation of duplicate classes as soon as the domain evolves (as an example, the allowed angles of travel could change or disappear entirely evolving this pathfinding problem into an any-angle pathfinding problem).

The actual [PathfindingAlgorithm](#)'s are modelled as abstract classes, this is because all pathfinding problems share the fact that they find a path between two points, and since most of the consumer classes of this module would require paths reaching multiple goals, we can place the shared functionality in this class. The fact this class is not concrete also allows us to decouple the consumer modules from a single concrete choice of [PathfindingAlgorithm](#).

The PathfindingGoal interface allows for any class with a `getCoordinates` method to become a target for pathfinding, this class is going to be implemented by each sensor. This is because the algorithms do not need to know anything about the objects to which we're passing other than their coordinates. Without this interface, the pathfinding module would

have to accept the [Sensor](#) class directly, which would couple these two modules and classes together. This would also defeat the purpose of abstracting pathfinding away.

The Obstacle interface allows for the inclusion of obstacles and consideration of obstacles in the pathfinding without coupling it to the problem domain.

The SpatialHash interface is modelled to allow different techniques of spatial hashing to be used with the same [PathfindingAlgorithm](#) without changing the class. The PathfindingHeuristic interface serves exactly the same role but for heuristic techniques.

Figure 2 is a summarised class diagram of the pathfinding module. Overall the class choices form an entirely decoupled module which allows for pathfinding in various planar domains employing, with the flexibility of mixing any combination of [PathfindingAlgorithm](#), PathfindingHeuristic and SpatialHash classes.

1.3 Simulation module

Once a mechanism for pathing between two points in space is found, the only problem left in creating a full path is finding out in which order to visit the sensors. This is the Travelling Salesman Problem, Many techniques exist for solving this problem and so general classes are established to allow a variety of methods to be created as needed. Figures 5 and 4 show a class diagrams of classes in the simulation module.

The CollectionOrderPlanner interface allows for the decoupling of specific implementations of TSP solvers from the consumers of this class, It would be possible to simply reference a specific implementation, but then swapping out the TSP solver would require entirely new classes, or modifications to the consumer class. An abstract [BaseCollectionOrderPlanner](#) is also included since all CollectionOrderPlanner's will use a distance matrix of some sort as well as a set of common optimisations which can be applied on top of any TSP solver (an example would be 2-opt). The PathPlanner interface is introduced for the same reason as above, to decouple the consumer of this class from any specific implementation. This enables flexibility in swapping out behaviour via different PathPlanner implementations. Since PathPlanners are decoupled from the specific pathfinding algorithm, these classes are tasked with the conversion of path points received from the [PathfindingAlgorithm](#), to [PathSegments](#) conforming to the design specification - a reading can only be performed at the end of a move. This conversion can be done in a variety of different ways, hence the interface - however a base abstract class is also provided to encapsulate the constraints of reading range and maximum amount of moves, should these constraints change in the specification, it would be easy to implement a new branch of these planners in the new domain.

Finally, [SensorDataCollector](#) is the interface representing the actual drone (or any data collector whatsoever), these classes apply the path and collection order planners to plan their journey, they may apply different strategies to how they use pathfinding and TSP solving to collect all the sensors, and prioritize the different goals differently. A base class is provided since all data collectors are expected to use both the planners, however should a collector be required which does not, this is also made possible thanks to this design.

CollectionOrderOptimiser is an interface for all general algorithms designed to improve already existing TSP tours. Examples of these include: 2-opt, k-opt, Genetic Algorithms, etc.. The interface allows for the application of a whole list of optimisers to an arbitrary path without knowing anything about the optimisers.

The [DistanceMatrix](#) abstract class encapsulates the general square adjacency matrix used to store distances between nodes in a graph, this allows for the swapping out of distance measures in the CollectionOrderPlanner's freely without changing the planners themselves. The Graph and [DistanceMatrix](#) classes completely define the domain of the problem together, and this makes this entire module applicable to a variety of different scenarios should new ones arise.

Figure 4 shows the relationship of this module to the pathfinding module. The abstract [Sensor](#) class implements the PathfindingGoal interface, and it is used directly by the PathPlanner's for pathfinding. The module defines its own specific [SearchNode](#) and Graph implementations which are propagated through the rest of the module. The [DirectedSearchNode](#) defines an additional field to store the direction from its parent, and the [ConstrainedTreeGraph](#) produces these nodes in agreement with the move length, direction, obstacle and boundary constraints. This effectively removes the need for any other class to deal with obstacle detection or working out any vector math, it's all encapsulated in this class. This centralisation of the collector's movement capability in a single class, means that any optimisations require only modifications to the graph. This also means that the entire array of classes in this module could suddenly be passed an entirely different graph within an entirely different domain, and they would all still work. This is rather excellent.

The [Building](#) class implements the Obstacle interface from the pathfinding module which allows it to be considered as an obstacle in pathfinding should need for it arise.

The [PathSegments](#) class is what encapsulates the movement of the collector, i.e. move in an integer direction followed by single sensor reading. While this could have been further abstracted, it is believed, that any form of movement can be approximated with this format, i.e. should a more complex path be necessary, the move length could be reduced and more angles allowed effectively increasing the flexibility of the drone extensively.

The Client module is what actually provides an implementation of the [Sensor](#) class, meaning that this module could work in entirely different scenarios, with readings in different ranges, sensors with more complex behaviours and so on.

1.4 Visualisation module

This module only needs to interface with the output classes of the simulation module, i.e. the [PathSegments](#) and Sensor classes but other than that, it is completely decoupled from the shape of the initial input data of the system. This is

uni-directional relationship is very desirable

Figure 3 shows the class diagram of the visualisation module.

The SensorCollectionVisualiser interface is one to again decouple the consumers of this module from any specific implementation, should the requirements for visualisation change, this can easily be done by creating a new visualiser class. The OutputFormatter class deals with writing the visualisations and flight paths to a file, these are made static as changes in output format are assumed to be very few in the future, and should such changes be required, new methods can be added to the formatter, but this is not a problem since changes to output format would likely require entirely different libraries and as such new classes dealing with those libraries.

[AQMapGenerator](#) is an example implementing class, which uses line strings and markers as in the specification. The usage of the AttributeMap interface allows for a lot of flexibility in the way the [AQMapGenerator](#) assigns colors and symbols, and should this behaviour need to be changed, it'd be very easy to do.

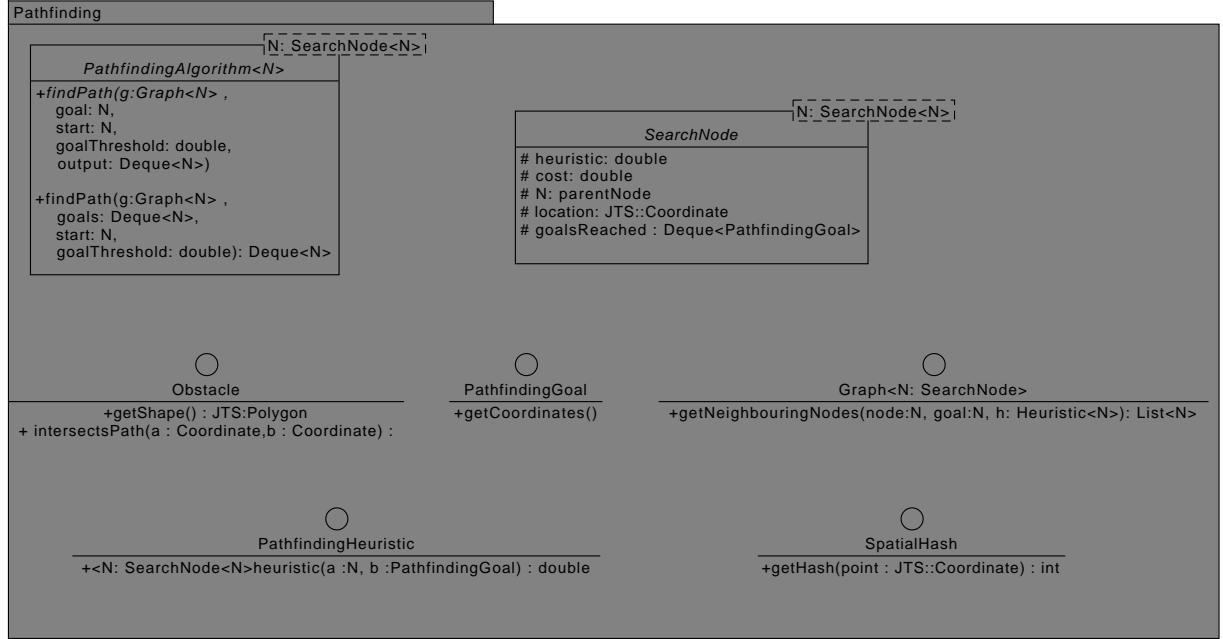


Figure 2: UML diagram of the pathfinding module

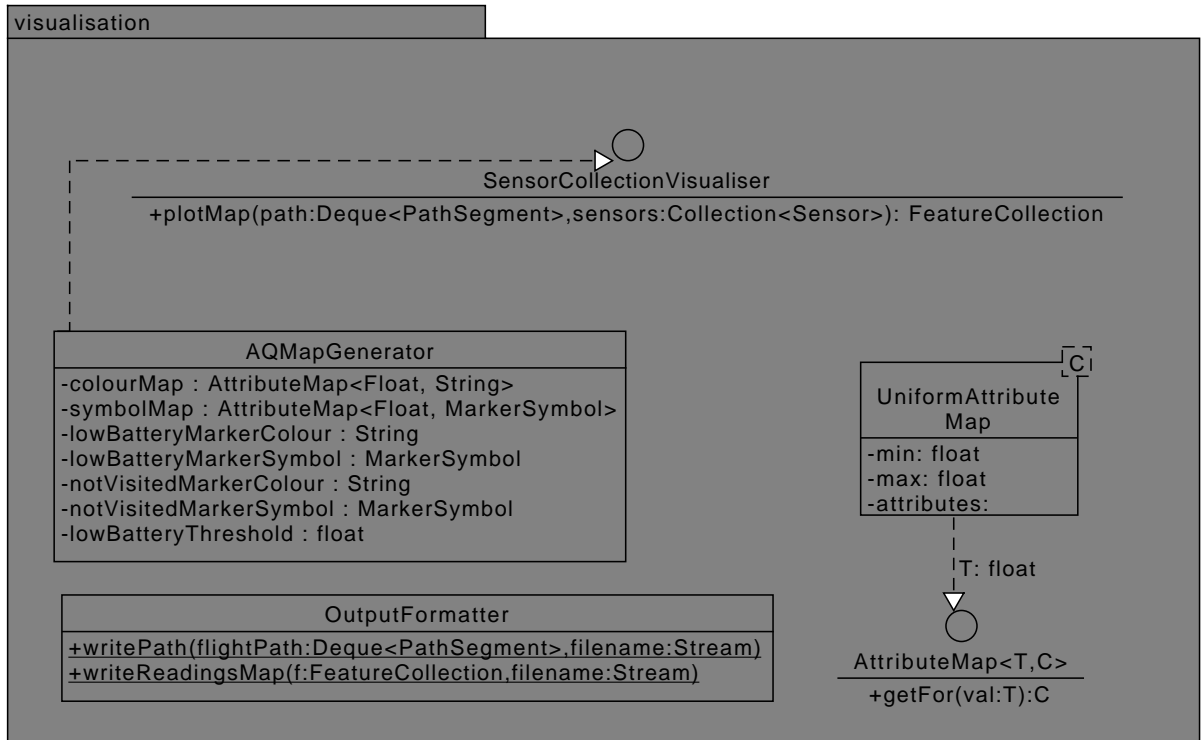


Figure 3: UML diagram of the Visualisation module

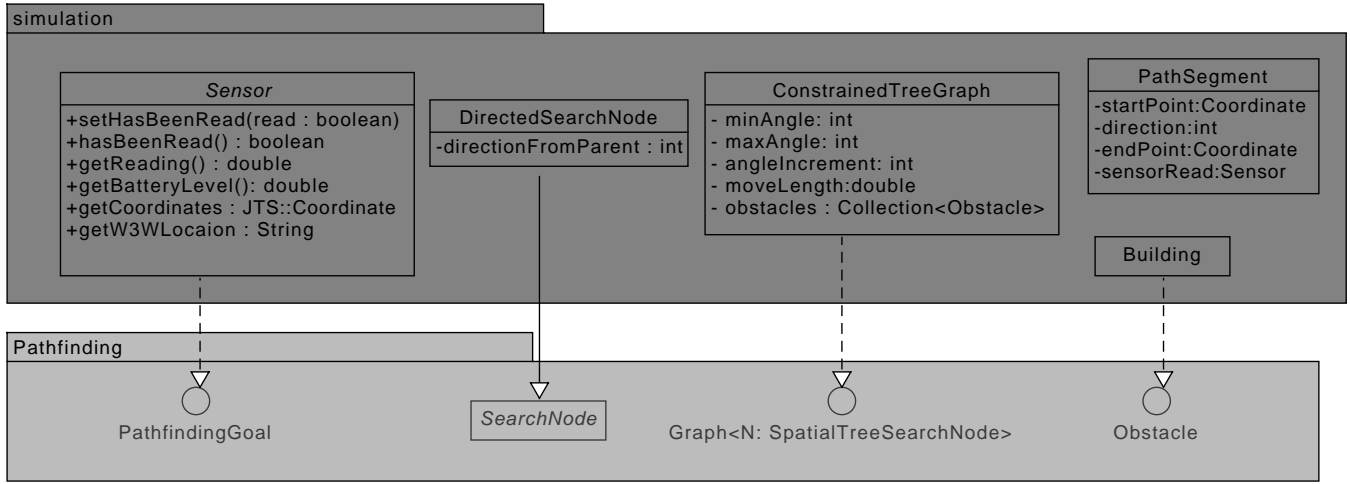


Figure 4: UML diagram of the Simulation module part 1

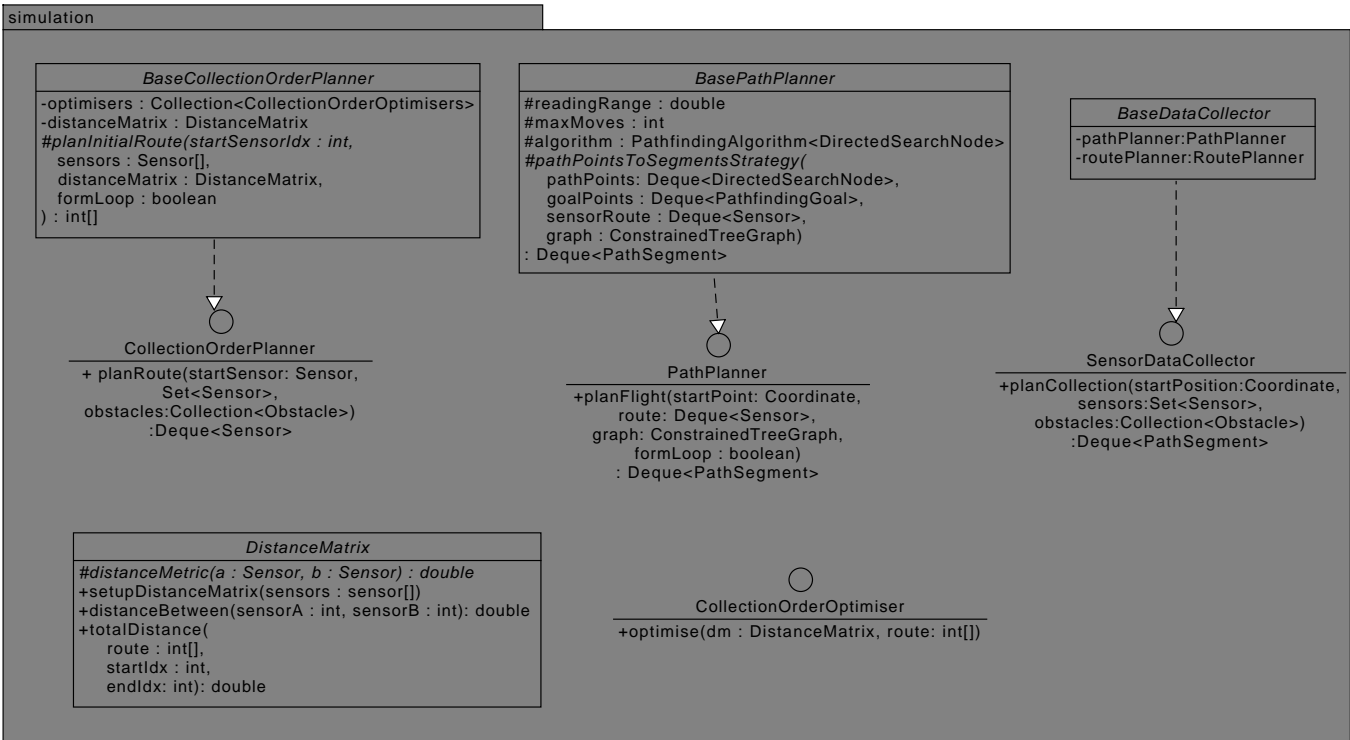


Figure 5: UML diagram of the Simulation module part 2

2 Drone Control Algorithm

2.1 TSP Solver

```
Algorithm NearestInsertion:
  Let R be the sensor nearest to the starting point
  Let U  $\leftarrow$  unvisited sensors
  Let T  $\leftarrow$  [R] be the current tour (implicitly looping)

  While U  $\neq \emptyset$ :
    R  $\leftarrow$  argmins $\in$ R distanceToTour(s)
    i  $\leftarrow$  argmini $\in$ I(T) insertionCost(i, R)
    T  $\leftarrow$  T with R inserted at i
    U  $\leftarrow$  U - {R}

  Return T

Function distanceToTour(s):
  minimum  $\leftarrow \infty$ 
  For t in T:
    If dist(s, t) < minimum:
      minimum  $\leftarrow$  dist(s, t)
  Return minimum

Function insertionCost(i, s):
  N  $\leftarrow$  T with s inserted at i
  Return euclidian length of N
```

Figure 6: Nearest Insertion TSP solver

The choice of which order to visit the sensors in is a general TSP problem.

The heuristic chosen to tackle this part of the problem was the [Nearest Insertion heuristic](#) followed by a number of [2-opt optimisations](#).

As pseudocode in Figure 6 shows, the heuristic inserts sensors into the tour one at a time, choosing the sensor which is closest to any sensor in the tour.

After a route is found, [2-opt optimisations](#) are applied to it to remove paths which are crossing each other. The 2-opt algorithm in each pass checks if reversing any sub-segment of all possible sub-segments of a full tour reduces its cost, and if so keeps the reversal. The passes are repeated until improvements fall under a threshold of 0.00003 degrees.

This algorithm turns out to be pareto-optimal [1] among a family of cutting-edge algorithms for problem sizes of between 30-50 vertices. Pareto-optimality means that this algorithm was either finding the best path compared to the other algorithms, or was finding one the quickest. Since this process yields the route built on the Minimum Spanning Tree, the path found will always be less than or equal to 2 times the optimal solution [2].

2.2 Generating neighbour nodes

All the other sections of the algorithm make use of the "neighbour" function, i.e. the function which generates neighbour nodes for any node representing a point on the map.

This is done using the [Bounding Volume Hierarchy](#) data structure. The structure allows for logarithmic time lookups of possibly colliding obstacles (with any shape). We do this by creating a binary tree whose nodes are defined by an Axis Aligned Bounding Volume enveloping all the obstacles present in the leaf nodes underneath the node. The root node then envelops all the obstacles present in the hierarchy.

When creating the tree we find the axis along which the difference between the extremal coordinates of the AABB's is the largest, i.e. the "longest axis". We then pick a splitting point on the axis and partition the shapes into the left and right sub trees according to which side of the splitting point they're on.

With this setup, the structure can tell us which obstacles are possibly colliding with any given shape by checking for collisions (cheaply) with the AABB's and only returning those leafs whose AABB's were collided with (possibly from both subtrees).

2.3 Pathfinding

```

Algorithm A*:
  Let R  $\leftarrow$  starting node
  let G  $\leftarrow$  goal coordinates
  Let O  $\leftarrow$  {R} // open set
  Let V  $\leftarrow$  {} // approximately visited

  While O  $\neq$   $\emptyset$ :
    R  $\leftarrow$   $\text{argmin}_{o \in O} \text{Fvalue}(o)$ 
    O  $\leftarrow$  O - R

    If isNearGoal(R):
      R.goalsReached  $\leftarrow$  R.goalsReached + G
      Return reconstructPath(R) // using parent references

    // generate neighbour nodes with appropriate costs and parent set to R
    // nodes colliding with obstacles or outside of boundary are excluded
    // this is done via Bounding Volume Hierarchies
    N  $\leftarrow$  neighbours(R)

    For n in N:
      hash  $\leftarrow$  cantorHash(n.x, n.y)

      If hash in V:
        Skip n

      V  $\leftarrow$  V + hash
      O  $\leftarrow$  O + n

  // no path found
  Return []

Function Fvalue(n):
  Return dist(n, G) * 1.5 + cost(n)

Function cantorHash(x, y):
  Let gridWidth  $\leftarrow$   $\frac{1}{75} * 0.0003$ 
  Let gridX, gridY  $\leftarrow$  coordinates of center of drone confinement area
  nx, ny  $\leftarrow$  (x - gridX) / gridWidth, (y - gridY) / gridWidth
  nx, ny  $\leftarrow$  makePositive( $\lfloor nx \rfloor$ ), makePositive( $\lfloor ny \rfloor$ )

  return  $\lfloor (((0.5 * (nx + ny)) * (nx + ny + 1)) + ny) \rfloor$ ;

Function makePositive(n):
  If n  $\geq$  0:
    Return 2n
  Otherwise:
    Return -2n - 1

```

Figure 7: Custom A* pathfinding algorithm

Pathfinding between any two points on the plane was carried out using the [A* tree search](#) algorithm using the euclidian distance heuristic. The algorithm is just a modified breadth-first search, where the search nodes are picked in order of least f value which is defined as: $f(n) = 1.5 \cdot h(n) + c(n)$ with $c(n)$ being the cost of reaching the node n from the start state, and $h(n)$ is the approximation of the cost of reaching the goal state from the node n.

The heuristic is relaxed by a factor of 1.5, to promote the expansion of straighter paths. This means that the paths are no longer optimal as the heuristic is no longer admissible, but since most paths would not be obstructed by buildings - this change does not actually impact the length of the paths all that much.

To speed up the algorithm [spatial hashing](#) was used to prune nodes which were visited (or visited close-enough), expressing their coordinates relative to a grid centered around the center of the boundary (with integer width square size), and hashing those coordinates using a modified cantor pairing allowing for negative values.

2.4 Path segmenting

Once we have a path of individual points, we need to convert it to a path of path segments with information about which sensor is read at each segment. We can do this by "sliding" a window over the points and always looking at two points at a time in order. We call the points at each position of the window P and N, where P is the point before N. The pathfinding algorithm also attaches a deque of sensors reached by each P/N. A naive sequential pairing does not work due to multiple goals allowed at each point, and the requirement of the drone to move before collecting the reading. We apply a number of rules to make a valid path (portrayed in Figure 8):

- While P attains a goal, create a proxy segment to any neighbouring node, and one back, assign the first sensor attained by P to the newest back-facing segment's N node
- If N attains no goals, see if the next segment's P node attains any sensors, if so "steal" one away from it
- while N attains more than one sensor reduce the number of attained sensors by creating proxy segments as above for each
- If N only attains one sensor or none, create P-N path segment as normal




Figure 8: Segmenting behaviours

2.5 Examples

The algorithm was tested on over 35000 configurations over the data provided. The hardest and easiest collection days' geojson visualisations are shown below (with the algorithm set to optimal parameters as given above)



Figure 9: geojson.io rendering of hardest collection at day 9-2-2020 with a starting point of -3.19087,55.945778, with 111 moves



Figure 10: geojson.io rendering of easiest collection at day 2-1-2021 with a starting point of -3.1878,55.9444, with 46 moves

The artifacts of path segmenting are evident in the harder day, where the drone wanders near closely spaced sensors on the bottom right it creates a horizontal proxy segment. But overall the algorithm manages very well, never carrying out the collection in more than 111 moves and averaging at $90 \text{ moves} \pm 6.7$. Execution times averaged at 103ms.

3 Class Documentation

uk.ac.ed.inf.aqmaps.visualisation

class AQMapGenerator

Visualiser of sensor data collections, portrays the flight path as a geojsononline string, and shows each sensor's reading using a combination of marker colour and symbol

Implements **SensorCollectionVisualiser**

Members

public AQMapGenerator (AttributeMap < Float , String > colourMap , AttributeMap < Float , MarkerSymbol > symbolMap)	Create an AQMapGenerator with the default low battery and non-visited symbol/colour values
public AQMapGenerator (AttributeMap < Float , String > colourMap , AttributeMap < Float , MarkerSymbol > symbolMap , String lowBatteryColour , MarkerSymbol lowBatterySymbol , String notVisitedColour , MarkerSymbol notVisitedSymbol)	Create an AQMapGenerator specifying all its parameters
public com.mapbox.geojson.FeatureCollection plotMap (Deque < PathSegment > flightPath , Collection < Sensor > sensorsToBeVisited)	Creates a FeatureCollection containing the line string of the flight path, and markers for each sensor given - styled according to its readings

enum MarkerSymbol

Available symbols for use with geojson, the toString methods are overridden to conform to the geojson marker list

class OutputFormatter

Formats and writes the flight path and geojson visualisation to a file

Members

public OutputFormatter ()	
public static void writePath (Deque < PathSegment > flightPath , OutputStream file)	Write flight path to given file
public static void writeReadingsMap (com.mapbox.geojson.FeatureCollection readingsMap , OutputStream file)	Write geojson readings visualisation to the given file

class UniformAttributeMap

A general attribute map which maps a range of values from a (min,max) range to attribute buckets of size (max-min)/buckets number uniformly.

Implements **AttributeMap**

Members

public UniformAttributeMap (Float min , Float max , C... attributes)	Create a new map with the given minimum and maximum range of values for input, values will be mapped to the list of attributes uniformly
public C getFor (Float value)	Retrieve attribute for the given input

class BVHNode

Bounding Volume Hierarchy Node. This class forms a tree of AABB (Axis aligned bounding boxes) for internal nodes and of any shapes at the leaf nodes. Allows for quick broad phase collision checks between objects. Will never return a false negative but might return false positives. I.e. this structure only tells you which objects are possibly colliding (whose AABB's intersect).

Extends **Shape**

Members

public BVHNode (Collection < T > shapes)	Construct a new bvh hierarchy with the given shapes at the leaf nodes
public Collection < T > getPossibleCollisions (org.locationtech.jts.geom.Geometry other)	Retrieves all possibly colliding objects with the geometry given.

class GeometryFactorySingleton

The geometry factory containing the precision model to be used when generating geometries with JTS

Members

public GeometryFactorySingleton ()	
public static org.locationtech.jts.geom.GeometryFactory getGeometryFactory ()	Retrieve the geometry factory, containing the precision model

class GeometryUtilities

A collection of utility geometry methods

Members

public GeometryUtilities ()	
public static org.locationtech.jts.geom.Coordinate MapboxPointToJTSCoordinate (com.mapbox.geojson.Point p)	Convert a mapbox point to a jts coordinate
public static com.mapbox.geojson.Point JTSCoordinateToMapboxPoint (org.locationtech.jts.geom.Coordinate p)	Convert a jts coordnate to a mapbox point
public static org.locationtech.jts.geom.Polygon MapboxPolygonToJTSPolygon (com.mapbox.geojson.Polygon p)	Convert mapbox polygon to jts polygon

class MathUtilities

Collection of general methods for dealing with angles and floating point comparisons

Members

public MathUtilities ()	
public static double angleFromEast (org.locationtech.jts.math.Vector2D a)	returns the angle from the eastern direction clockwise between 0 and 360 of the vector
public static double oppositeAngleFromEast (double angle)	Return the angle representing the given angle incremented by 180 degrees (but witin 360 degrees)
public static org.locationtech.jts.math.Vector2D getHeadingVector (double angle)	gets unit vector in the direction of angle
public static boolean thresholdEquals (double a , double b , double epsilon)	Returns true if the given values are within a threshold of each other
public static boolean thresholdEquals (double a , double b)	Returns true if the given values are within a threshold of each other. Uses default epsilon
public static boolean thresholdEquals (org.locationtech.jts.geom.Coordinate a , org.locationtech.jts.geom.Coordinate b , double epsilon)	Returns true if the given coordinates are within a threshold of each other
public static boolean thresholdEquals (org.locationtech.jts.geom.Coordinate a , org.locationtech.jts.geom.Coordinate b)	Returns true if the given coordinates are within a threshold of each other. uses default epsilon

abstract class BasePathPlanner

Base class for planners with a limited number of maximum moves and a minimum reading range. All inheriting path planners must make sure that each sensor is read at the endpoint of some path segment and that that sensor is within reading range of the endpoint. They also must make sure that the path is under the maximum move limit

Implements **PathPlanner**

Members

public BasePathPlanner (double readingRange , int maxMoves , PathfindingAlgorithm < DirectedSearchNode > algorithm)	Create a new path planner with the given constraints, and using the given pathfinding algorithm
public Deque < PathSegment > planPath (org.locationtech.jts.geom.Coordinate startCoordinate , Deque < Sensor > route , ConstrainedTreeGraph graph , boolean formLoop)	Plans the exact path required to reach all the given sensors, the specific constraints on placed on the route are decided by the specific implementation of the planner itself. Plans the path starting from the given start coordinate, reaching all the given sensors and forming a loop back if the form loop argument is given
protected abstract Deque < PathSegment > pathPointsToSegmentsStrategy (Deque < DirectedSearchNode > pathPoints , Deque < PathfindingGoal > goalsRoute , Deque < Sensor > sensorRoute , ConstrainedTreeGraph graph)	The main defining characteristic of a constrained path planner. Converts a path of points to a path of path segments needs to make sure that each pathfinding goal is visited only in the end segment of some path segment in range. the passed deque arguments will be consumed

class PathSegment

represents a singular move made by the sensor data collector
each move follows the pattern of: move and then collect reading,
we cannot collect a reading in a move unless we have moved

Members

```
public PathSegment (  
    org.locationtech.jts.geom.Coordinate startPoint ,  
    int direction ,  
    org.locationtech.jts.geom.Coordinate endPoint ,  
    Sensor sensorRead )
```

Creates a path segment from the start and end points, the direction of movement and the sensor read if any

class SimplePathPlanner

This path planner will apply some simple optimisations in order to produce a smaller number of path segments than the naive implementation.

Implements **PathPlanner** Extends **BasePathPlanner**

Members

```
public SimplePathPlanner (  
    double readingRange ,  
    int maxMoves ,  
    PathfindingAlgorithm<DirectedSearchNode> algorithm )
```

```
protected Deque<PathSegment> pathPointsToSegmentsStrategy (  
    Deque<DirectedSearchNode> pathPoints ,  
    Deque<PathfindingGoal> goalsRoute ,  
    Deque<Sensor> sensorRoute ,  
    ConstrainedTreeGraph graph )
```

The main defining characteristic of a constrained path planner. Converts a path of points to a path of path segments
needs to make sure that each pathfinding goal is visited only in the end segment of some path segment in range.
the passed deque arguments will be consumed. This planner will try to perform some simple optimisations in order to shorten the route.
In order to produce a valid route this planner will introduce proxy segments which go back and forth between the nearest neighbour whenever a sensor is read at the start point of a segment or if more than one sensor is read at the endpoint. The optimisations currently include:

1) if the current segment does not read anything at the end point and the next reads a sensor at its start point, we "absorb" that sensor into the current segment.

uk.ac.ed.inf.aqmaps.simulation.planning.collectionOrder.optimisers

class Optimiser2Opt

An optimiser which performs the 2-opt algorithm to remove crossings in the path

Implements **CollectionOrderOptimiser**

Members

```
public Optimiser2Opt ( double epsilon )
```

Construct a 2 opt optimiser with the given epsilon threshold. The threshold determines the minimum decrease in path cost required for the optimiser to keep optimising each loop.

```
public void optimise (  
    DistanceMatrix distanceMatrix ,  
    int[] path )
```

uk.ac.ed.inf.aqmaps.simulation.planning.collectionOrder

abstract class BaseCollectionOrderPlanner

Collection order planners generate good traversal orders between the given set of sensors, where "good" criteria are defined by each implementation of the collection planner.

Members

```
public BaseCollectionOrderPlanner (  
    Collection<CollectionOrderOptimiser> optimisers ,  
    DistanceMatrix distMat )
```

Creates a collection order planner with the given optimisers and distance matrix method

```
public Deque<Sensor> planRoute (  
    Sensor startSensor ,  
    Set<Sensor> sensors ,  
    boolean formLoop )
```

Generates a collection order over the sensors.

```
protected abstract int[] planInitialRoute (  
    int startSensorIdx ,  
    Sensor[] sensors ,  
    DistanceMatrix distanceMatrix ,  
    boolean formLoop )
```

Drafts a route between the given sensors, using the given matrix. If form loop is true then the route will also begin and end on the same sensor

class GreedyCollectionOrderPlanner

Plans a collection of sensor data in a greedy order
i.e. by picking the closest sensor at each step.

Extends [BaseCollectionOrderPlanner](#)

Members

public GreedyCollectionOrderPlanner (Collection < CollectionOrderOptimiser > optimiser , DistanceMatrix distMat)	Creates a new Greedy planner with the given optimisers and distance matrix
protected int[] planInitialRoute (int startSensorIdx , Sensor [] sensors , DistanceMatrix distanceMatrix , boolean formLoop)	Drafts a route between the given sensors, using the given matrix. If form loop is true then the route will also begin and end on the same sensor

class NearestInsertionCollectionOrderPlanner

Collection order planner which employs the nearest insertion method to try and pick the best route.

Extends [BaseCollectionOrderPlanner](#)

Members

public NearestInsertionCollectionOrderPlanner (Collection < CollectionOrderOptimiser > optimisers , DistanceMatrix distMat)	Constructs a new NI planner with the given optimisers and distance matrix
protected int[] planInitialRoute (int startSensorIdx , Sensor [] sensors , DistanceMatrix distanceMatrix , boolean formLoop)	Drafts a route between the given sensors, using the given matrix. If form loop is true then the route will also begin and end on the same sensor

uk.ac.ed.inf.aqmaps.simulation.planning

class ConstrainedTreeGraph

A graph which imposes angle, move length and boundary (+ obstacle) constraints for the nodes, and does not keep track of already produced nodes (tree search) i.e. a new node is returned each time.
The angle system needs to allow for each possible angle to have a "complement angle" which takes you back to where you started if you moved in its direction after stepping in any possible angle.
the min and max angle need to cover a range of 360 degrees - the angle increment .

Implements [SearchGraph](#)

Members

public ConstrainedTreeGraph (int minAngle , int maxAngle , int angleIncrement , double moveLength , Collection < Obstacle > obstacles , org.locationtech.jts.geom.Polygon boundary)	Constructs new tree graph with the given angle and move constraints and obstacles, and the bounds of the zone outside of which coordinates are not produced
public List < DirectedSearchNode > getNeighbouringNodes (DirectedSearchNode node)	Returns the neighbours of the given node within the graph
public int getClosestValidAngle (double direction)	returns the closest valid angle in this graph to the given angle

abstract class DistanceMatrix

Class which stores distance information between sensors

Members

public DistanceMatrix ()	Creates new blank distance matrix
public void setupDistanceMatrix (Sensor [] sensors)	Fills in the distance matrix with distance data for the given sensors, each sensor is referred to by its index in the sensors array given later.
public double distanceBetween (int a , int b)	returns the distance from sensor a to sensor b at the given indices in the sensor list
public double totalDistance (int[] route , int startIdx , int endIdx)	calculates the total distance of the subpath specified with start and end indices within the given path of sensor indices.
public double totalDistance (int [] route)	Calculates the total length of the given path (not including looping back)
protected abstract double distanceMetric (Sensor a , Sensor b)	the specific distance measure used to calculate distances. Does NOT have to be symmetric

class EuclidianDistanceMatrix

Distance matrix using euclidian distance as the value for distances

Extends [DistanceMatrix](#)

Members

public EuclidianDistanceMatrix ()	
protected double distanceMetric (Sensor a , Sensor b)	the specific distance measure used to calculate distances. Does NOT have to be symmetric

class GreatestAvoidanceDistanceMatrix

Distance matrix using the greatest avoidance distance as the distance metric. This distance is calculated by forming a minimum bounding circle around all obstacles between any two sensors and calculating the length of the path which "wraps" around the circle

Extends [DistanceMatrix](#)

Members

public (Collection < Obstacle > obstacles)	initialize blank distance matrix with the given obstacles
protected double distanceMetric (Sensor a , Sensor b)	the specific distance measure used to calculate distances. Does NOT have to be symmetric

uk.ac.ed.inf.aqmaps.simulation.collection

abstract class BaseDataCollector

Each data collector follows the same pattern, it uses a path planner to find a way between two points, as well as a collection order planner which sets out the route around all the sensors. Each collector may use this data differently, for example it may discard the path given by a path planner under certain circumstances, or change the route mid-way.

Implements [SensorDataCollector](#)

Members

public BaseDataCollector (PathPlanner fp , BaseCollectionOrderPlanner rp)	Constructs a data collector with the given path planner and the given collection order planner
---	--

class Drone

the drone collector is not constrained by the map layout, if the graph (or map) allows a node to be reached the drone can fly through it, the graph itself may impose constraints indirectly, but the drone assumes absolutely no restrictions in its movements.

Implements [SensorDataCollector](#) Extends [BaseDataCollector](#)

Members

public Drone (PathPlanner fp , BaseCollectionOrderPlanner rp)	
public Deque< PathSegment > planCollection (org.locationtech.jts.geom.Coordinate startCoordinate , Set < Sensor > sensors , ConstrainedTreeGraph graph , boolean formLoop , int randomSeed)	Plans the best order of visiting sensors (best being defined by the collector itself) and creates a detailed ordered path segment collection which when followed will allow a successful collection. The sensors returned will have their state set to either read or not read, depending on whether they are reached within the path or not

uk.ac.ed.inf.aqmaps.simulation

class Building

An obstacle with a polygonal shape representing a building

Implements [Obstacle](#) Implements [Shape](#)

Members

public Building (org.locationtech.jts.geom.Polygon shape)	initialize new building with the given shape
public boolean intersectsPath (org.locationtech.jts.geom.Coordinate a , org.locationtech.jts.geom.Coordinate b)	Returns true if the line formed from a to b intersects this obstacle

class DirectedSearchNode

A data structure representing a tree search node for spatial pathfinding problems with integer angles.

Extends [SearchNode](#)

Members

public DirectedSearchNode (org.locationtech.jts.geom.Coordinate location , DirectedSearchNode parent , int directionFromParent , double heuristic , double cost)	Creates fully specified tree search node
---	--

abstract class Sensor

Sensors contain a reading value and a battery level as well as the what 3 words location and coordinates of the sensor. This sensor will contain placeholder values for readings and battery status until it is set to have been read.

Implements **PathfindingGoal**

Members

public Sensor (org.locationtech.jts.geom.Coordinate coordinates , float reading , float batteryLevel , String W3WLocation)	
public void setHaveBeenRead (boolean read)	Sets the sensor state to "read". If the sensor is not set to have been read, asking it for readings and battery levels will return placeholder values
public float getReading ()	Return the reading at this sensor. Will return NaN until the setHaveBeenRead(true) has been called
public float getBatteryLevel ()	Return the battery level at this sensor. Will return NaN until the setHaveBeenRead(true) has been called
public boolean hasBeenRead ()	Returns true if this sensor has been read (i.e. setHaveBeenRead(true) has been called on this object)

class SensorDataCollectorFactory

A utility class for instantiating sensor data collectors with correct parameter values in one call

Members

public SensorDataCollectorFactory ()	
public static SensorDataCollector createCollector (ConstrainedTreeGraph g , double readingRange , int maxMoves , CollectorType collectorType , PathfindingHeuristicType heuristicType , CollectionOrderPlannerType plannerType , DistanceMatrixType matrixType)	Create a new collector within the given domain and with the given parameters, filling in missing parameters with optimal values.

enum SensorDataCollectorFactory.CollectionOrderPlannerType

The type of the collection order planner to use

enum SensorDataCollectorFactory.CollectorType

The type of sensor data collector to use

enum SensorDataCollectorFactory.DistanceMatrixType

The type of distance matrix to use for the TSP solver

enum SensorDataCollectorFactory.PathfindingHeuristicType

The type of pathfinding heuristic to use

uk.ac.ed.inf.aqmaps.pathfinding.heuristics

class StraightLineDistance

The simplest heuristic, uses the euclidian distance between the node and its goal as the value of its heuristic.

Implements **PathfindingHeuristic**

Members

public StraightLineDistance ()	Initializes a new instance of the straight line heuristic with a relaxation factor of 1 (i.e. no relaxation)
public StraightLineDistance (double relaxationFactor)	Initializes a new instance of the straight line heuristic with the given relaxation factor. Some pathfinding algorithms such as Astar will run much faster with a relaxed heuristic, but the relaxed solution might not be optimal (path cost <= relaxationFactor * optimal path cost)
public <T extends SearchNode<T>> double heuristic (T a , PathfindingGoal b)	Returns the straight line distance between the nodes multiplied by the relaxation factor

uk.ac.ed.inf.aqmaps.pathfinding.hashing

class GridSnappingSpatialHash

hashes real coordinates by scaling them and "snapping" them to a grid..

Implements **SpatialHash**

Members

public GridSnappingSpatialHash (double gridSize , org.locationtech.jts.geom.Coordinate gridCenter)	Create new instance of grid snapping hash
public int getHash (org.locationtech.jts.geom.Coordinate a)	Returns a hash for the given coordinate. A good hash will be equal for points which are near each other according to some metric, and different for ones that are not.

class PointGoal

A point goal is the simplest kind of goal, a single location in space.

Implements [PathfindingGoal](#)

Members

public PointGoal (org.locationtech.jts.geom.Coordinate goal)	Creates a new point goal
--	--------------------------

class AstarTreeSearch

Classic pathfinding algorithm, modified BFS which uses both the cost to reach a node and the predicted cost from that node to the goal to chose the nodes to be expanded next. This version of Astar treats the search as a tree search and so uses a hashing function to determine if a node has been visited yet.

Extends [SearchNode](#)

Members

public AstarTreeSearch (PathfindingHeuristic heuristic , SpatialHash hash)	Creates a new instance of Astar search with the given heuristic and spatial hashing function.
public void findPath (SearchGraph <T> g , PathfindingGoal goal , T start , double goalThreshold , Deque <T> output)	Description copied from class: PathfindingAlgorithm

abstract class PathfindingAlgorithm

Pathfinding algorithms operate over any valid search nodes and graph definitions. The graph defines the transition function between one node and its neighbours while the search nodes are used as the path constructing object. Any number of pathfinding algorithms can be defines in these terms, both tree and graph searches are possible with the correct set of graph and node definitions.

Extends [SearchNode](#)

Members

public PathfindingAlgorithm ()	
public abstract void findPath (SearchGraph <T> g , PathfindingGoal goal , T start , double goalThreshold , Deque <T> output)	Finds a path from the start to the goal node and outputs the result into the provided deque. If a path doesn't exist no nodes will be added to the output, if it does at least one node will be added. The nodes will have their goalsReached deque's set to the corresponding goals that can be reached from their locations within the goal threshold. One node can reach multiple nodes
public Deque <T> findPath (SearchGraph <T> g , Deque < PathfindingGoal > route , T start , double goalThreshold)	Finds a path from the start node through the provided route. if at any point there exists no path between the given goals, the route will halt before that segment (so might be empty).
protected boolean isAtGoal (double threshold , T node , PathfindingGoal goal)	Checks whether the given node is within a threshold away from the goal
protected void reconstructPathUpToIncluding (T node , Deque <T> out , T limitNode)	reconstructs the path to node and deposits it in the given queue

class SearchNode

Search Nodes are used to hold the path information in pathfinding algorithms including heuristic, cost and parent node values. Search nodes also contain information about which pathfinding goals can be achieved from their position (can be multiple).

Members

public SearchNode (org.locationtech.jts.geom.Coordinate location , T parent , double cost)	Creates a new search node which is fully specified apart from the heuristic value
public SearchNode (org.locationtech.jts.geom.Coordinate location , T parent , double heuristic , double cost)	Creates a fully specified search node

class SensorData

Object representing the reading, battery and location data of a sensor

Members

public SensorData ()	Empty constructor, necessary for proper de-serialization
public SensorData (String W3WLocation , float battery , float reading)	Constructs a new sensor data object

class W3WAddressData

An object containing information about a what-3-word square including it's coordinates and meta-data

Members

public W3WAddressData ()	Empty constructor, necessary for proper de-serialization
public W3WAddressData (String country , W3WSquareData square , String nearestPlace , com.mapbox.geojson.Point coordinates , String words , String language , String map)	Create a new W3WAddressData object from the given address information

class W3WSquareData

Object representing information about a what-3-words square's coordinates.

Members

public W3WSquareData ()	Empty constructor, necessary for proper de-serialization
public W3WSquareData (com.mapbox.geojson.Point southwest , com.mapbox.geojson.Point northeast)	Create a new W3WSquareData object from the diagonal corners of a square

class AQSensor

An implementation of the [Sensor](#) interface for consumption in the simulation module. Conveniently can be formed from sensor and address data corresponding to the data received via API client.

Implements [PathfindingGoal](#) Extends [Sensor](#)

Members

public AQSensor (SensorData sensorData , W3WAddressData w3WAddressData)	Construct new AQ sensor from sensor data representing the sensor's readings and battery information, and with the address data containing the w3w address words and coordinates of the sensor
---	---

class AQWebServerClient

The web server client communicates with a server which contains the necessary information and retrieves it

Implements [ClientService](#)

Members

public AQWebServerClient (HttpClient client , URI APIBaseURI)	Create a new web server client, if the base api uri is localhost then the base api would be : http://localhost
public List< SensorData > fetchSensorsForDate (LocalDate date)	Retrieves all SensorData for sensors to be collected on the given date
public W3WAddressData fetchW3WAddress (String wordAddress)	Retrieves detail information about a what-3-words address (i.e. word.word.word)
public W3WAddressData fetchW3WAddress (String w1 , String w2 , String w3)	Convenience method, accepts a split w3w address
public com.mapbox.geojson.FeatureCollection fetchBuildings ()	Retrieves all the no fly-zones i.e. the buildings present

References

- [1] The Metric Travelling Salesman Problem: The Experiment on Pareto-optimal Algorithms, Sergey Avdoshin, E.N.Beresneva 2017.
- [2] https://aswani.ieor.berkeley.edu/teaching/FA13/151/lecture_notes/ieor151 lec17.pdf, The Traveling Salesman Problem Professor Z. Max Shen